

Segmentation Introduction

J. Cory – Embedded Vision Specialist



ML Task: Segmentation

> Task of semantic segmentation:

- >> Input: Image
- >> Output: Image
 - >> Label each pixel in the image with a category label
 - >> Don't differentiate instances, only care about pixels
- >> Evaluation metric: mIOU




Deep CNN Architecture

Encoder
Downsample
(feature extractor!)

- > VGG
- > Inception
- > ResNet

Decoder
Upsample

- > Deconv
- > Interpolation




Semantic Segmentation

> Applications of semantic segmentation:

- >> Lane detection, drivable space detection, environmental perception in automotive.

> Task of instance segmentation:

- >> Input: Image
- >> Output: Image
 - >> Label different individuals of the same class
 - >> Don't care about background
- >> Evaluation metric: AP, average precision




Instance Segmentation

Segmentation Output Data

> Segmentation Networks Output

- For these kind of networks, the output layer is typically Crop and the size of the output layer will be same as the input size of the network.
- The output will be written sequentially in raster scan order inside output buffers.
- Each pixel represent the class ID/Label.



- Example:
For AlexNet-FCN, the last layer is Crop and the output size is same as the input size i.e. 500 x 500.

Metrics

> In general mean IoU is the metric of choice for segmentation

>> MS COCO uses the following metrics:

Leaf category metrics:

mIoU	% Mean IoU (primary challenge metric)
fIoU	% Frequency Weighted IoU
mAcc	% Mean Accuracy
pAcc	% Pixel Accuracy

Supercategory metrics:

mIoU ^S	% Mean IoU
fIoU ^S	% Frequency Weighted IoU
mAcc ^S	% Mean Accuracy
pAcc ^S	% Pixel Accuracy

1. Intersection Over Union (IoU) is computed as $\text{IoU} = \text{TP} / (\text{TP} + \text{FP} + \text{FN})$, where TP=true positives, FP=false positives, and FN=false negatives.
2. Mean IoU (mIoU) is the average over the IoUs of each class.
3. Frequency Weighted IoU (fIoU) weights every class proportional to the number of pixels for that class.
4. Mean Accuracy (mAcc) denotes the fraction of correct pixels per class averaged over classes.
5. Pixel Accuracy (pAcc) denotes the fraction of correct pixels.
6. Note that metrics are computed on the *pixels* in each image, *not on object instances*.
7. All metrics are computed only on valid pixels (91 stuff classes + 1 other class). Unlabeled pixels are not considered.
8. If there are no ground-truth pixels for a class in the validation/test set, that class is removed from consideration.

>> CityScapes and Kitti Datasets use the following metrics:

- IoU class: Intersection over Union for each class $\text{IoU} = \text{TP} / (\text{TP} + \text{FP} + \text{FN})$
- iloU class: Instance Intersection over Union $\text{iloU} = \text{iTP} / (\text{iTP} + \text{FP} + \text{iFN})$
- IoU category: Intersection over Union for each category $\text{IoU} = \text{TP} / (\text{TP} + \text{FP} + \text{FN})$
- iloU category: Instance Intersection over Union for each category $\text{iloU} = \text{iTP} / (\text{iTP} + \text{FP} + \text{iFN})$

Semantic Segmentation General Approaches

> Fully Convolutional Networks (FCN)

- >> Uses deconvolutions to expand back up to original image size
- >> Later versions also add skip connections to maintain some of input resolution

> SegNet

- >> Similar to FCN but instead of short cut connections for features, indices from maxpooling are copied (makes it more memory efficient than FCN)

> Dilated Convolutions

- >> Intended to increase resolution/FoV without increasing number of weights

> DeepLab (One variant holds the leaderboard on PASCAL VOC)

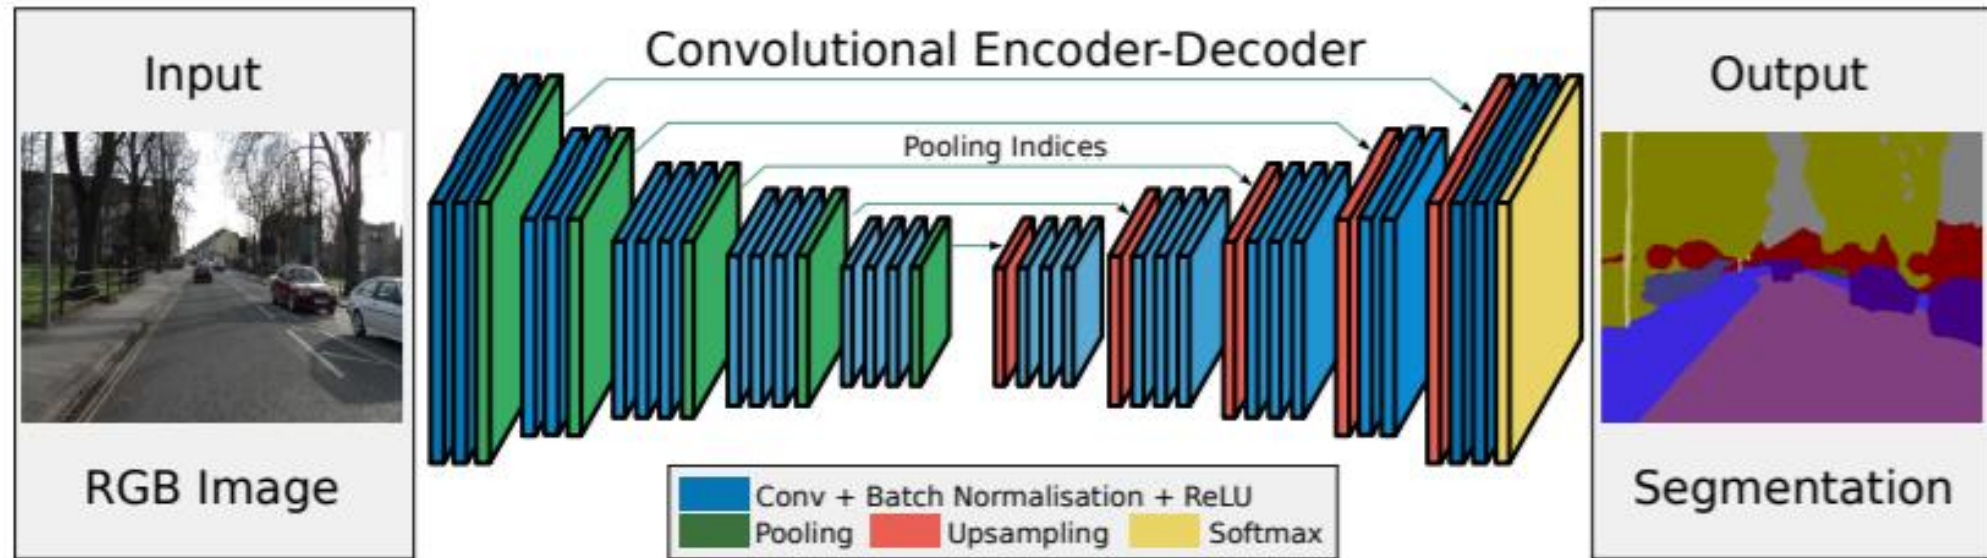
- >> Starts with Dilated convolutions and adds conditional receptive fields
- >> Also incorporates multiscale (image pyramid approach)

> FPN (ResNeXt152-FPN was the winning entry for MS COCO Stuff - 2017)

- >> Creates a pyramid of feature maps during feature extraction at $\frac{1}{2}$ scale for each layer, then upsamples back to original image size and provides skip connections from input maps to output


Source: <http://blog.qure.ai/notes/semantic-segmentation-deep-learning-review#segnet>

Segnet (Encoder/Decoder) Approach

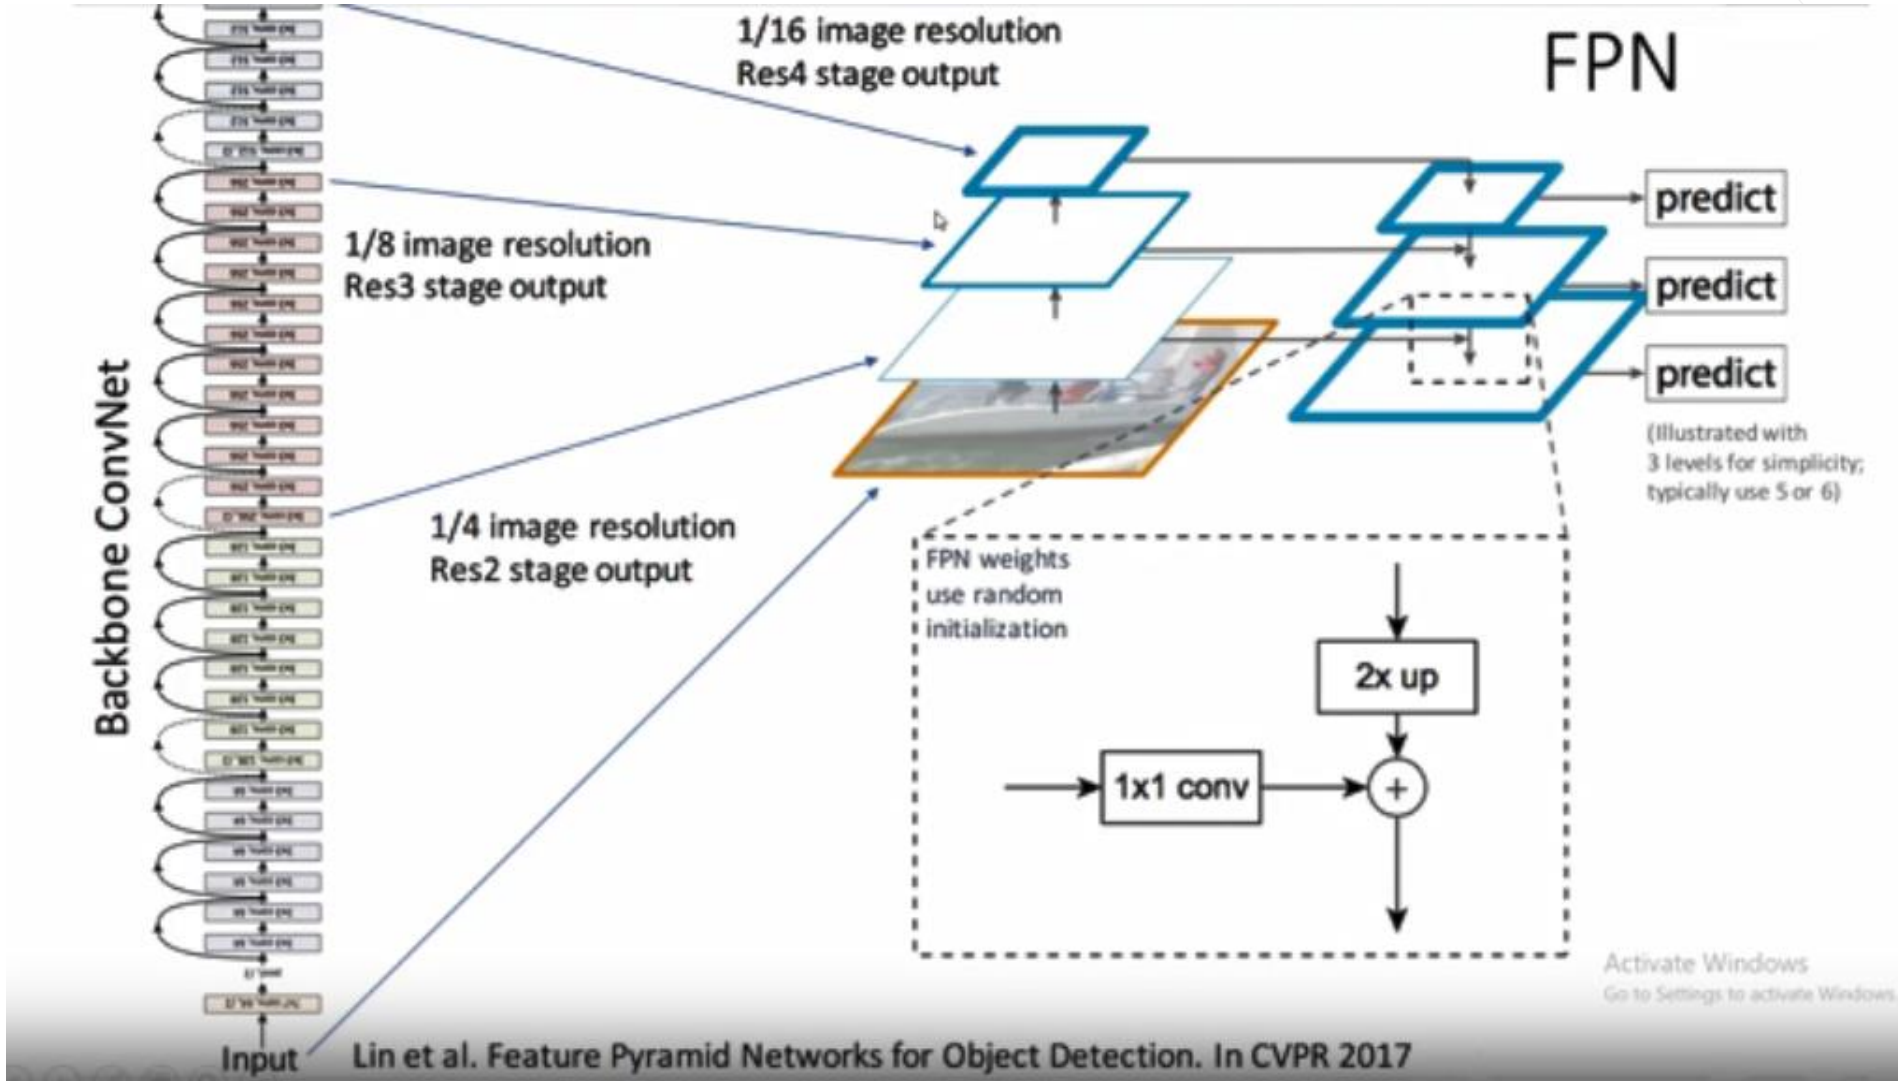


- > There are different flavors of Segnet
- > First 13 convolutional layers correspond to VGG16 (encoder stage)
- > Decode stage consists of 13 layers which correspond to the encode layers
 - >> For pass through paths, convolution is performed with max pooling indices and trainable decoder filters (instead of deconvolution)

Feature Pyramid Networks



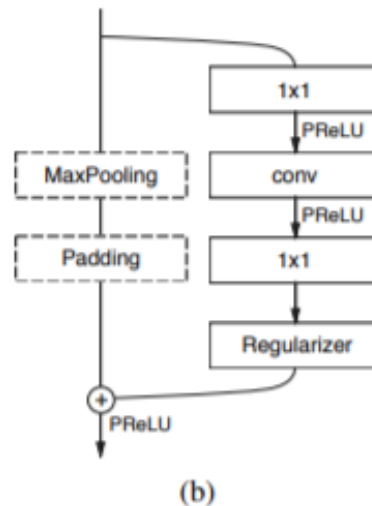

FPN Example with Resnet backbone



ENet Overview

> ENet builds on the encoder/decoder network approach similar to Segnet

- >> In general, it implements figure a for the initial stage, then figure b for the bottleneck stages as are shown in the table
- >> The latter stages of ENet (decoder stage) expand back up to the original image size
 - Pooling is also used to forward layers from the encoder through to the decoder stage to help maintain spatial coherence



Name	Type	Output size
initial		$16 \times 256 \times 256$
bottleneck1.0	downsampling	$64 \times 128 \times 128$
4× bottleneck1.x		$64 \times 128 \times 128$
bottleneck2.0	downsampling	$128 \times 64 \times 64$
bottleneck2.1		$128 \times 64 \times 64$
bottleneck2.2	dilated 2	$128 \times 64 \times 64$
bottleneck2.3	asymmetric 5	$128 \times 64 \times 64$
bottleneck2.4	dilated 4	$128 \times 64 \times 64$
bottleneck2.5		$128 \times 64 \times 64$
bottleneck2.6	dilated 8	$128 \times 64 \times 64$
bottleneck2.7	asymmetric 5	$128 \times 64 \times 64$
bottleneck2.8	dilated 16	$128 \times 64 \times 64$
<i>Repeat section 2, without bottleneck2.0</i>		
bottleneck4.0	upsampling	$64 \times 128 \times 128$
bottleneck4.1		$64 \times 128 \times 128$
bottleneck4.2		$64 \times 128 \times 128$
bottleneck5.0	upsampling	$16 \times 256 \times 256$
bottleneck5.1		$16 \times 256 \times 256$
fullconv		$C \times 512 \times 512$

Reference: <https://arxiv.org/pdf/1606.02147.pdf>

ESPNet Overview (1)

> ESPNet (Efficient Spatial Pyramid Network) has several flavors

- >> Types a-c produce reduced size feature map outputs and type d produces a segmentation map that is the same size as the input image
- >> The flavors implemented in the following tutorial are ESPNet-C (c) and the full ESPNet (d)
 - These two types are a sort of feature pyramid type of network
 - Pass through concepts of Segnet are also applied in types b, c, and d

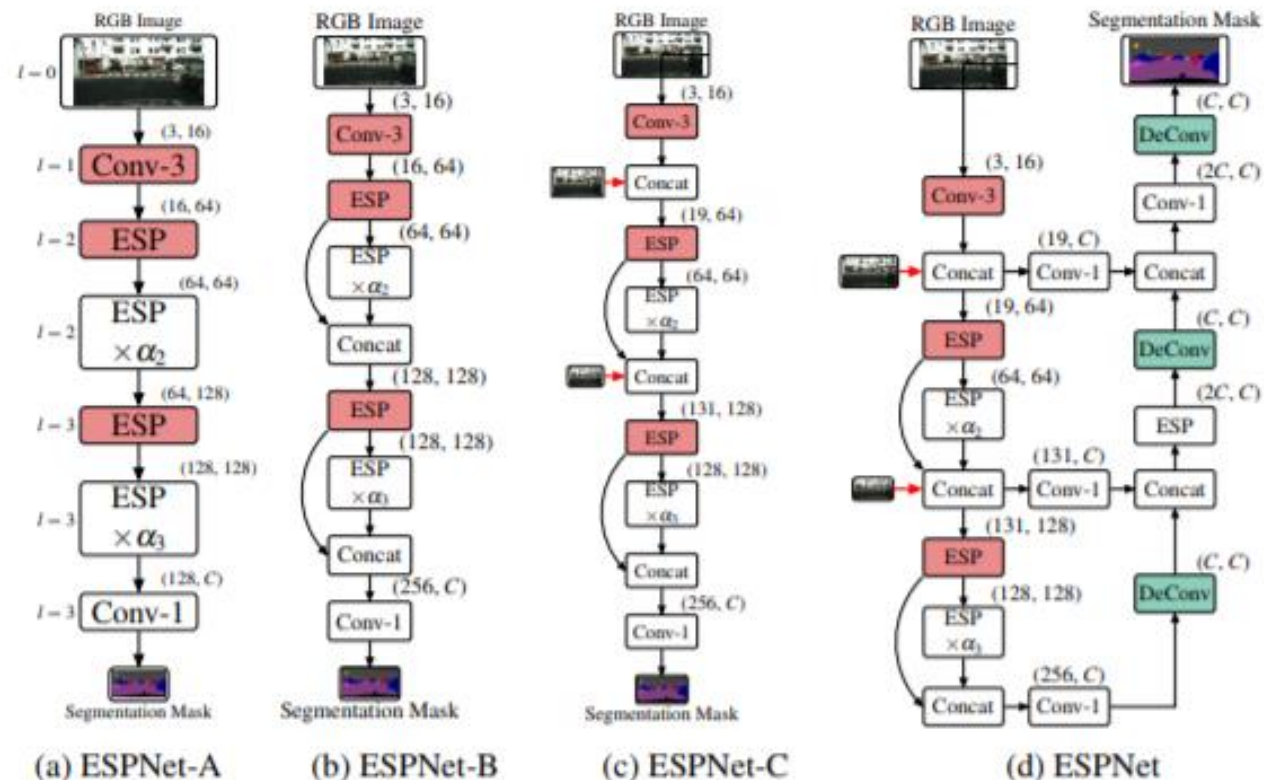
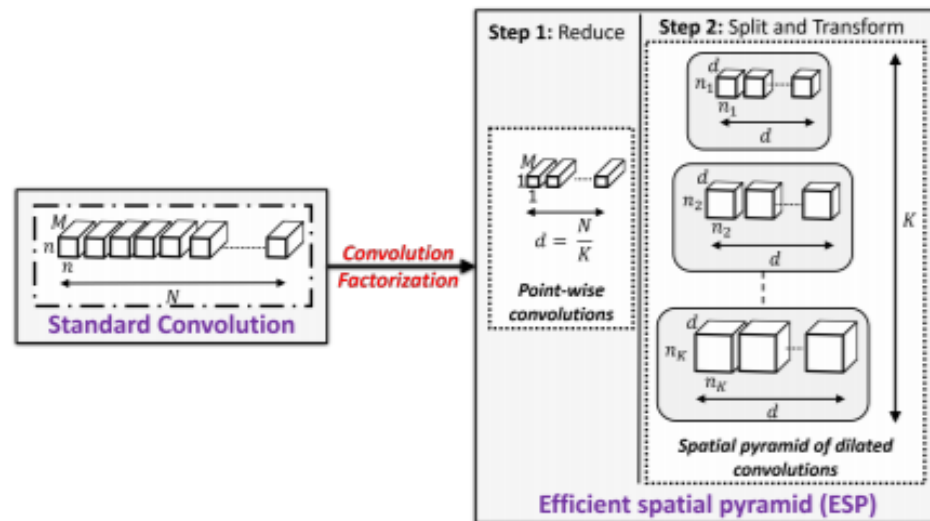


Fig. 4: The path from ESPNet-A to ESPNet. Red and green color boxes represent the modules responsible for down-sampling and up-sampling operations, respectively. Spatial-level l is indicated on the left of every module in (a). We denote each module as (# input channels, # output channels). Here, Conv- n represents $n \times n$ convolution.

ESPNet Overview (2)

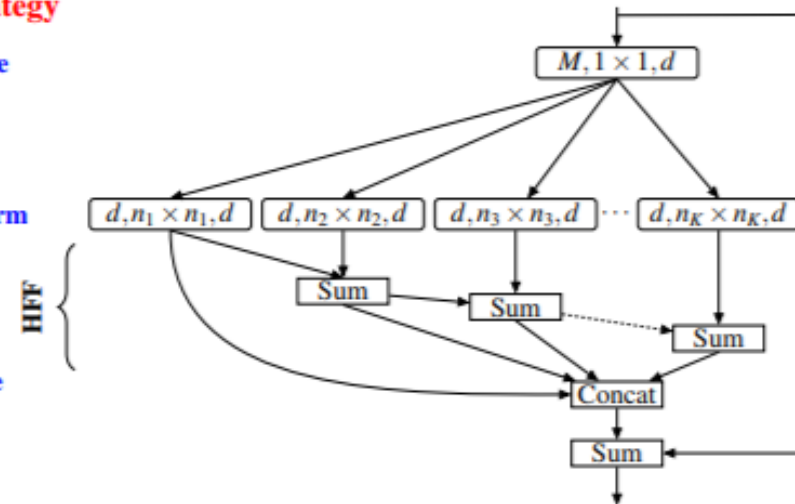
- > One of the novel features of ESPNet are ESP modules which are a factorized form of convolutions
 - >> Factorization is a way to optimize the convolution operation by decomposing it into a spatial pyramid of point-wise and dilated convolutions
 - >> This reduces the computational complexity of the convolution operation



(a)

ESP Strategy

Reduce
Split
Transform
Merge



(b)

Reference: <https://arxiv.org/pdf/1803.06815.pdf>

Adaptable.
Intelligent.

